Discretize mean-field limit using JKO, see if it is similar to GD.

1 CURRENT IN CODE

From one dimensional data, we add a dimension filled with ones to act as a bias for the first layer. The output of one ReLU neuron for one data point $(x, 1) \in \mathbb{R}^2$:

$$w, b, \alpha \in \mathbb{R} \rightarrow \max(0, wx + b)\alpha$$

The loss against labels $y_j \in \mathbb{R}$ using squared loss of the whole network of neurons is the double sum:

$$\mathcal{L} = \sum_{j=1}^{n} \left(\left(\sum_{i=1}^{m} \max(0, w_i x_j + b_i)\alpha_i \right) - y_j \right)^2$$

We can define a wasserstein gradient flow on the mean-field limit of this network, this requires taking an infinite-width ReLU network where parameters are described by a measure μ, and its output by an integral:

$$\int_{\mathbb{R}^2} m((w, b); x) \mathrm{d}\mu((w, b))$$

To simplify things, we restrict α_i to $\{-1, 1\}$ and to not be a trainable parameter anymore. We keep the same expressivity(as long as we provide both a positive($\alpha_i = 1$) and negative($\alpha_i = -1$) version of the neuron) but slightly alter the training dynamic in some cases. For example , we can match the output of one neuron (of the original network) by simply scaling the first layer by the seconder layer (α):

$$\max(0, w_i x + b_i)\alpha_i = \max\left(0, |\alpha_i|(w_i x + b_i)\right) \mathrm{sign}(\alpha_i)$$

Our network with restricted α_i would describe this neuron using only two trainable parameters: $(|alpha_i|w_i, |\alpha_i|b_i)$ and fix its sign in the output.

The measure is on the parameter space. In order to do simulations we discretize the parameter space, by taking a uniform grid in $\mathbb{R}^2$ centered on $(0, 0)$: $(w_i, b_i)_{i=1,\dots m}$

We can see that we have the same output and expressivity as the regular ReLU network by taking a measure $\mu = \sum_{i=1}^{m} p_i \delta_{\theta_i = w_i}$ with $(\sum_i p_i = 1)$ and $m((w_i, b_i); x) = \max(0, w_i x + b_i)\alpha_i$, we have this equality:

$$\int_{\mathbb{R}^2} m((w, b); x) \mathrm{d}\mu((w, b)) = \sum_{i=1}^{m} \max(0, w_i x_j + b_i)\alpha_i p_i$$

In this case, the first layer is fixed: the change of direction($\frac{-b_i}{w_i}$) and slope (w_i) of a neuron is described by a mass displacement from point A to point B.

The movemement is described by a PDE and simulated on a grid. Each point i of the grid has a weight $p_i \in \mathbb{R}$, and as a whole $p \in \mathbb{R}^m$ is the discretized distribution.

The same wasserstein gradient flow can be computed by this step:

$$\mu(t+1) = \underset{\mu \in \mathcal{M}(\Theta)}{\arg\min} F(\mu) + \frac{1}{2\gamma} W_2(\mu; \mu(t))$$

However, directly computing the wasserstein distance is too hard. We will compute the entropic wasserstein flow:

$$p_{t+1} := \text{Prox}_{\tau f}^{W_\gamma}(p_t)$$

$$= \underset{p \in \text{simplex}}{\arg\min}\ W_\gamma(p, q) + \tau f(p)$$

$$= \underset{p \in \text{simplex}}{\arg\min} \left(\min_{\pi \in \Pi(p,q)} \langle c, \pi \rangle + \gamma E(\pi) \right) + \tau f(p)$$

Where π is a mapping, c the ground cost for every point on the grid. When the ground cost between two points in the euclidian space is $c_{i,j} = \|x_i - x_j\|^2$, (and $\gamma = 0$, f smooth...), this scheme formaly discretize the above mentionned PDE.

To do the step above, we'll use a bregman splitting approach that replace the single implicit W_γ proximal step by many iterative KL implicit proximal steps. Specifically(?) Dykstra's algorithm for JKO stepping. This involve using the gibbs kernel:$\xi = e^{-\frac{c}{\gamma}} \in \mathbb{R}_{+,*}^{N \times N}$

Algorithm 1 JKOstep

1: $p \leftarrow p_0 \in \mathbb{R}^m$
2: $q_{\text{norm}} \leftarrow \|p\|^2$
3: $a, b \leftarrow \mathbf{1}, \mathbf{1} \in \mathbb{R}^m$ ▷ Initialize vectors with ones
4: **for** $i \leftarrow 1$ **to** T **do**
5:     $p \leftarrow \text{prox}_{\tau/\gamma}^{\text{KL}}(\xi b)$
6:     $a \leftarrow p/(\xi b)$
7:     $\text{ConstrEven} \leftarrow \frac{\|b \cdot (\xi a) - q\|}{q_{\text{norm}}}$
8:     $b \leftarrow q/(\xi a)$
9:     $\text{ConstrOdd} \leftarrow \frac{\|a \cdot (\xi b) - p\|}{q_{\text{norm}}}$
10:     **if** $\text{ConstrOdd} < \text{tol}$ **and** $\text{ConstrEven} < \text{tol}$ **then**
11:         **break**
12:     **end if**
13: **end for**

---

2   Classic setup

Data $x_j \in \mathbb{R}^d$ and labels $y_j \in \mathbb{R}$, $j = 1, .., n$

First layer $w_i \in \mathbb{R}^d$, second layer $\alpha_i \in \mathbb{R}$, $i = 1, .., m$

$\gamma > 0$ step-size, β regularization

$$\mathcal{L}(W, \alpha) = \sum_{j=1}^{n} \Big(\underbrace{\sum_{i=1}^{m} \max(0, w_i^\top x_j) \alpha_i}_{\text{Network's Output}} - y_j \Big)^2 + \lambda \underbrace{\sum_{i=1}^{m} \|w_i\|_2^2 + \alpha_i^2}_{\text{Weight Decay}}$$

Discret time.

Full-batch gradient descent

$$(W, \alpha)_{t+1} = (W, \alpha)_t - \gamma \nabla \mathcal{L}((W, \alpha)_t)$$

Implicit

$$\theta_{t+1} = \underset{\theta}{\arg\min}\ \mathcal{L}(\theta) + \frac{1}{2\gamma} \|\theta - \theta_t\|$$

Continuous time.

Taking $\gamma \to 0$, we get the gradient flow: $\frac{\mathrm{d}\theta_t}{\mathrm{d}t} = -\nabla\mathcal{L}(\theta_t)$. We make ReLU differentiable with $\sigma'(0) = 0$ as justified in (Boursier et al.).

3 Infinite width, using a measure: mean-field

Mean-field limit(Chizat & Bach): For a sufficiently large width, the training dynamics of a NN can be coupled with the evolution of a probability distribution described by a PDE.

If [...] converges, with $m \to \infty$ (many-particle limit), our particles of interest converges to a Wasserstein gradient flow of F:

$$\partial\mu_t = -\mathrm{div}(v_t\mu_t) \text{ where } v_t \in -\partial F'(\mu_t)$$

$$\int_\Theta m(\theta;x)\mathrm{d}\mu(\theta) = \frac{1}{m}\sum_{i=1}^m \langle w_i, x_j\rangle_+ \alpha_i$$

Different ways to use a measure to represent the neurons of a two layer network:

- $\Theta = \mathbb{R}^d \times \mathbb{R}$, measure $\mu = \frac{1}{m}\sum_{i=1}^m \delta_{\theta_i=(w_i,\alpha_i)}$, output of one neuron $m(\theta = (w,\alpha);x) = \langle x, w\rangle_+ \alpha$: (works, output matches discrete)
- $\Theta = \mathbb{R}^d$, measure $\mu = \frac{1}{m}\sum_{i=1}^m \alpha_i \delta_{\theta_i=w_i}$ output of one neuron $m(\theta = w;x) = \langle x, w\rangle_+$ (works)
- $\Theta = \mathbb{R}^d \times \mathbb{R}^d$, output of one neuron $m(\tilde{w}_+, \tilde{w}_-, x) = \langle \tilde{w}_+, x\rangle - \langle \tilde{w}_-, x\rangle$ (works, separate neg and positive)
- $\Theta = (S^{d-1} \times \mathbb{R})$, output of one neuron $m((d,\tilde{\alpha});x) = \tilde{\alpha}\langle d, x\rangle = \tilde{\alpha} \; \mathbb{1}_{\langle d,x\rangle>0}$ (works), mapping: $d = \frac{w}{\|w\|}$ and $\tilde{\alpha} = \|w\|\alpha$. Gradient are not equal to discrete.

3.1 Algorithm, discretize the measure's space

Take a grid of N points in Θ, we can match the notation above by taking a neuron for each point of the grid $m = N$.

$$\mu(t+1) = \underset{\mu\in\mathcal{M}(\Theta)}{\arg\min} F(\mu) + \frac{1}{2\gamma}W_2(\mu;\mu(t))$$

3.2 JKO

What we compute by using the entropic JKO flow iterations.

$$\forall t > 0, p_{t+1} := \mathrm{Prox}_{\tau f}^{W_\gamma}(p_t)$$
$$= \underset{p\in\text{simplex}}{\arg\min} W_\gamma(p,q) + \tau f(p)$$
$$= \underset{p\in\text{simplex}}{\arg\min}\left(\min_{\pi\in\Pi(p,q)} \langle c,\pi\rangle + \gamma E(\pi)\right) + \tau f(p)$$

- Meta Optimal Transport (paper) and (code git): InputConvexNN to predict solution of OT problem
- JKOnet (paper) and (code git):
  - /models -> sinkhorn loss defined in loss.py, differentiable loop in fixed point.py
  - next step: trying to create the right Geometry object from OTT library, which is what's used for sinkhorn

3.3 Papers

The algo we try to implement

Paper with a specific case that doesn't match ours:

In the future, large-scale waserstein gradient flows

3.3.1 Grid problems

The grid currently dictate the neuron's scale, giving multiple choices. One solution: duplicate each neuron, make one with a small scale and one with a very big scale.

References

Etienne Boursier, Loucas Pillaud-Vivien, and Nicolas Flammarion. Gradient flow dynamics of shallow relu networks for square loss and orthogonal inputs. URL http://arxiv.org/abs/2206.00939.

Lénaïc Chizat and Francis R. Bach. On the global convergence of gradient descent for over-parameterized models using optimal transport. In Samy Bengio, Hanna M. Wallach, Hugo Larochelle, Kristen Grauman, Nicolò Cesa-Bianchi, and Roman Garnett (eds.), *Adv. Neural Inf. Process. Syst. 31 Annu. Conf. Neural Inf. Process. Syst. 2018 NeurIPS 2018 Dec. 3-8 2018 Montr. Can.*, pp. 3040–3050. URL https://proceedings.neurips.cc/paper/2018/hash/a1afc58c6ca9540d057299ec3016d726-Abstract.html.